

# PARALLEL LOOPS



AARHUS  
UNIVERSITY  
DEPARTMENT OF ELECTRICAL AND COMPUTER  
ENGINEERING

SWD



# RECAP LAMBDA EXPRESSIONS

```
Func<int, int> square = x => x * x;
```

```
Console.WriteLine(square(5)); _____> ?  
Console.WriteLine(square);    _____> ?
```

```
Action<int> print = i => Console.WriteLine(i);
```

```
print(5);  
print(10);
```

```
Func<int> longComputation = () =>  
{  
    int i;  
    for (i = 0; i < 10; i++)  
    {  
        // Do some work  
        Thread.Sleep(1);  
    }  
    return i;  
};  
Console.WriteLine(LambdaAsParameters.MeasureTime(longComputation));
```

```
2 references  
internal class LambdaAsParameters  
{  
    2 references | 2/2 passing  
    public static long MeasureTime(Func<int> someFun)  
    {  
        Stopwatch stopwatch = new Stopwatch();  
        stopwatch.Start();  
        Console.WriteLine(someFun());  
        stopwatch.Stop();  
        return stopwatch.ElapsedMilliseconds;  
    }  
}
```



# RECAP LAMBDA EXPRESSIONS

```
int stop = 5;
Func<int> longComputation = () =>
{
    int i;
    for (i = 0; i < stop; i++)
    {
        // Do some work
        Thread.Sleep(1);
    }
    return i;
};
Console.WriteLine(LambdaAsParameters.MeasureTime(longComputation));

int stop = 5;
Func<int> longComputation = () =>
{
    int i;
    for (i = 0; i < stop; i++)
    {
        // Do some work
        Thread.Sleep(1);
    }
    return i;
};
stop = 10;
Console.WriteLine(LambdaAsParameters.MeasureTime(longComputation));
```

2 references

```
internal class LambdaAsParameters
{
    2 references | 🟢 2/2 passing
    public static long MeasureTime(Func<int> someFun)
    {
        Stopwatch stopwatch = new Stopwatch();
        stopwatch.Start();
        Console.WriteLine(someFun());
        stopwatch.Stop();
        return stopwatch.ElapsedMilliseconds;
    }
}
```



# LOOPING IN APPLICATION

---

- A significant part of application work is done in loop constructs
- Often, the loop iterations are independent of each other
- When iterations are indeed independent, they may be executed in parallel (PoPP: “delightfully parallel execution”)

```
for(int i=0; i<10; i++)  
{  
    WriteLine("i is " + i)  
}
```



```
Parallel.For(0, 10, i =>  
{  
    Console.WriteLine("i is " + i);  
});
```

# AN INITIAL IMPLEMENTATION OF PARRALLEL LOOPS

” just to appreciate the problems of  
partitioning  
- TROELS FEDDER

# PLANNING IMPLEMENTATION

---

- The parallel loop signature:

```
public static void MyParallelFor(  
    int inclusiveLowerBound,  
    int exclusiveUpperBound,  
    Action<int> body);
```

- Partitioning to individual threads (“1 thread per core”)

```
int size = exclusiveUpperBound - inclusiveLowerBound;  
int numProcs = Environment.ProcessorCount;  
int range = size / numProcs;
```

Example: size = 35, numProcs = 4  $\Rightarrow$  range = 8



# THE IMPLEMENTATION

---


```
public static void MyParallelFor(int inclusiveLowerBound, int exclusiveUpperBound,
Action<int> body) {
    // Determine size of each partition of work (size/nCores) - static partitioning
    int size = exclusiveUpperBound - inclusiveLowerBound;
    int numProcs = Environment.ProcessorCount;
    int range = size / numProcs;

    // Initialize threads to do work
    var threads = new List<Thread>(numProcs);
    for (int p = 0; p < numProcs; p++)
    {
        int start = p * range + inclusiveLowerBound;
        int end = (p == numProcs - 1) ? exclusiveUpperBound : start + range;
        threads.Add(new Thread(() => {
            for (int i = start; i < end; i++) body(i);
        }));
    }

    // Start and await threads
    foreach (var thread in threads) thread.Start(); // Start them all
    foreach (var thread in threads) thread.Join(); // wait on all
}
```

# SO ARE ALL GOOD?


- The cost of creating/killing threads is massive!
  - 1 MB stack, 100.000-200.000 cycles for construction/teardown
- The danger of oversubscription:
  - `MyParallelFor()` may itself be called in parallel → 8 (, 12, 16, ...) threads for 4 CPUs
  - OS spends time context-switching (takes time, kills caches)
  - Oversubscription example: “yellow is pain!”






# STATIC PARTITIONING

- Static partitioning – load imbalance
  - The in-equivalent workload for each iteration → some threads complete before others
  - Threads represent a static partitioning – threads cannot “help each other out”




# PARTITIONING IS IMPORTANT!

- Assume a parallel loop over  $N = [1; 12]$ , where iteration  $i$  takes  $i$  seconds
- Total time to complete =  $1+2+3+...+11+12 = 78$  secs
- Ideal load balance (dual core):  $78:2 = 39$  secs to complete



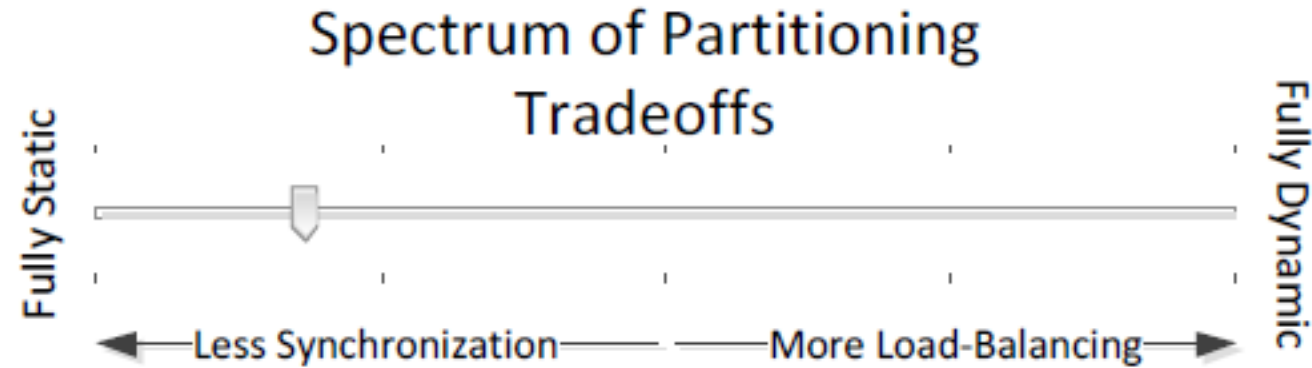
- Our loop uses static load balancing
  - Thread 1: iteration 1 thru 6 21 secs
  - Thread 2: iteration 7 thru 12 57 secs

Show finish order, not proportional time



- Total time to complete loops: 57 secs (46% longer than ideal)

# STATIC OR DYNAMIC PARTITIONING



- Effective static partitioning requires a priori knowledge of execution time, but requires no synchronization
  - But is less-than-ideal
- Effective *dynamic* partitioning requires no knowledge but requires synchronization
  - Or the threads will step on each other's toes

# ENTER PARALLEL EXTENSIONS

---

- In .NET 4, the class `Parallel` was introduced
- `Parallel` provides 3 static methods (and overloads)
  - `Parallel.For(start, end, Action)`
  - `Parallel.ForEach(collection, Action)`
  - `Parallel.Invoke(Action)`
- `Parallel.For()/ForEach()` provides a number of benefits
  - Exception handling, thread-local state, nested parallelism, dynamic thread count, and sophisticated load balancing

In short: The works!

# Parallel.For()

---

- Parallel version of a regular for-loop:

```
public static ParallelLoopResult For(  
    int fromInclusive, int toExclusive, Action<int> body);
```

## Example

```
for (int i = 0; i < nCalculations; i++)  
    C[i] = Math.Sqrt(Math.Pow(A[i], 2.0) + Math.Pow(B[i], 2.0));
```

Translate to

```
Parallel.For(0, nCalculations, i =>  
    {  
        C[i] = Math.Sqrt(Math.Pow(A[i], 2.0) + Math.Pow(B[i], 2.0));  
    }  
);
```



# Parallel.ForEach()

---

- Parallel version of iteration over collection (foreach)

```
public static ParallelLoopResult ForEach<TSource>(
    IEnumerable<TSource> source, Action<TSource> body);
```

## Example

```
foreach (var arg in feArgs) {
    arg.C = Math.Sqrt(Math.Pow(arg.A, 2.0) + Math.Pow(arg.B, 2.0));
}
```

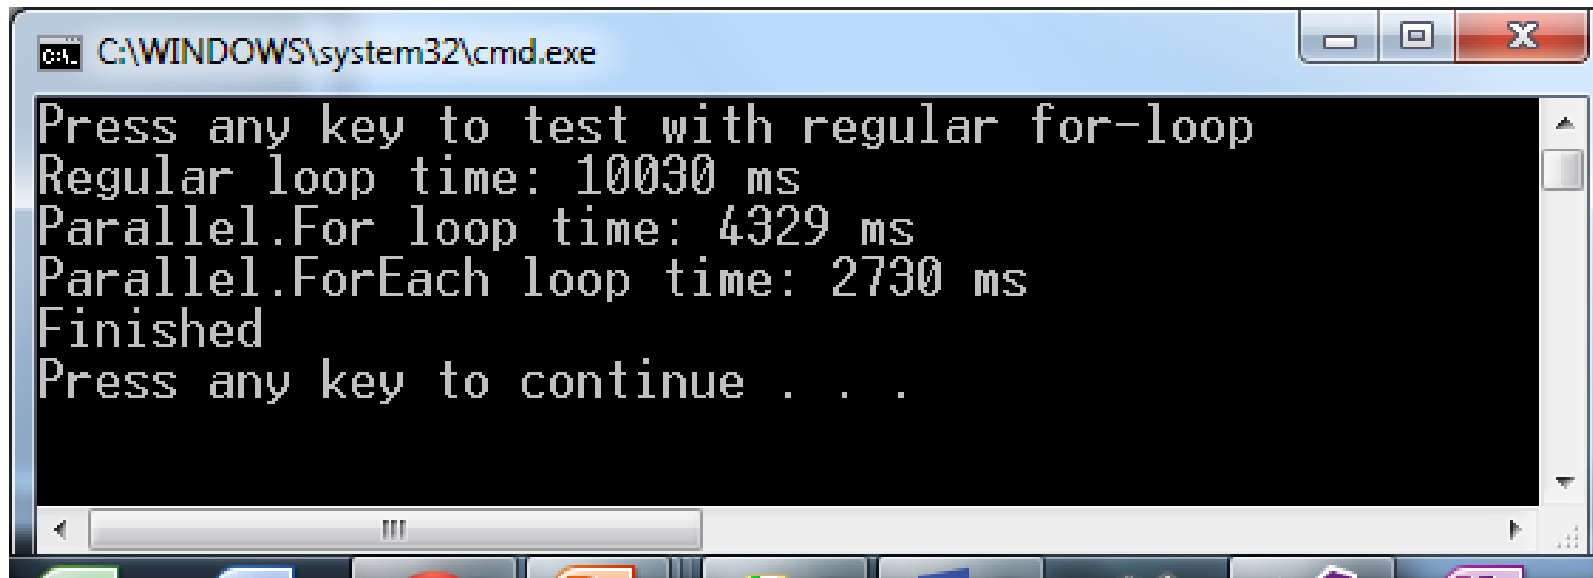
Translate to

```
Parallel.ForEach(feArgs, arg => {
    arg.C = Math.Sqrt(Math.Pow(arg.A, 2.0) + Math.Pow(arg.B, 2.0));
});
```



# EXAMPLE: DISTANCE CALCULATIONS

- 1.000.000 coordinates, calculate the distance to each other
- Using regular `for(...)`, `Parallel.For()`, and `Parallel.ForEach()`



```
C:\WINDOWS\system32\cmd.exe

Press any key to test with regular for-loop
Regular loop time: 10030 ms
Parallel.For loop time: 4329 ms
Parallel.ForEach loop time: 2730 ms
Finished
Press any key to continue . . .
```

# THE DANGER ZONES

---

- To use delightfully parallel loops, the iterations *must* be independent

```
Parallel.For(2, nCalculations, i =>  
    {  
        a[i] = a[i-1] + a[i-2]; // Oh God, the pain...the PAIN!!  
    }  
);
```

- Iterations are not always [0..n)
  - Downward iterations `for(..; ..; i--)`
  - Stepped iterations `for(..; ..; I += 2)`
- Very small loop bodies may defeat parallelisation
  - Overhead in delegate invocation and load balancing synchronization



# OTHER LANGUAGES

---

- Java
  - `list.parallelStream().forEach()`
  - `IntStream.range(0,10)`  
`.parallel()`  
`.forEach(i -> ...);`
- C++
  - OpenMP
  - AMP - Accelerated Massive Parallelism  
Using the GPUs on the graphics card.
- Python
  - multiprocessing or joblib



AARHUS  
UNIVERSITY