

PARALLEL TASKS IN C#/.NET



AGENDA

- Task Basics
- The Default Task Scheduler
- Passing data to asynchronous work
- The Task-Based Asynchronous Pattern (TAP)

TASK BASICS

- A task is an isolated, logical unit of work – a sequential operation fit for **parallelization**
 - A task is NOT a thread
 - A task is a sequential operation
 - `System.Threading.Tasks`
- Also: Use parallel tasks when you have several distinct **asynchronous** operations (see later under TAP)
- Tasks are provided in the Task Parallel Library (TPL)
 - Part of Microsoft Parallel Extensions for .NET (along with PLINQ)
 - TPL dynamically scales the degree of parallelism to most efficiently use all processors available.
 - TPL assists in the partitioning of work and scheduling of tasks in the .NET thread pool.

STARTING TASKS

Task parallelism at its simplest:

```
// Good ole sequential code
public void DoAll()
{
    DoLeft();
    DoRight();
}
```

```
// Using Parallel.Invoke()
public void DoAll()
{
    // Parallel.Invoke() automatically creates tasks
    // for the arguments, then awaits for completion
    Parallel.Invoke(DoLeft, DoRight);
}
```

```
// Using Task.Run() - from .NET 4.5
public void DoAll()
{
    Task t1 = Task.Run((Action) DoLeft);
    Task t2 = Task.Run((Action) DoRight);
    Task.WaitAll(t1, t2); // Wait for both tasks to
    complete
}
```

```
// Using TaskFactory
public void DoAll()
{
    Task t1 = Task.Factory.StartNew(DoLeft);
    Task t2 = Task.Factory.StartNew(DoRight);
    Task.WaitAll(t1, t2); // Wait for both tasks to
    complete
}
```

Tasks do not necessarily begin executing on creation. They are placed in a *work queue* from which a *task scheduler* remove and schedule them for execution when e.g. a core is available

WAITING FOR TASK COMPLETION

```
// Using WaitAll() or WaitAny()
public void DoAllUsingWait()
{
    Task t1 = Task.Run((Action)DoLeft);
    Task t2 = Task.Run((Action)DoRight);

    Task.Wait(t1);
    Task.Wait(t2);

// --- OR ---

    Task.WaitAll(t1, t2); // Wait for both tasks to complete

// --- OR ---


    Task.WaitAny(t1, t2); // Wait for the first task to complete
}
```

Tasks defer exception handling. The caller of `Wait*()` gets the exception. This allows “sequential-style” exception handling

THE DEFAULT TASK SCHEDULER


- TPL uses a task scheduler to schedule the tasks
 - Default task scheduler
 - Custom task scheduler (not covered)
- Tasks are generally very fine-grained work items
- TPL uses worker threads to execute tasks
 - Worker threads are managed by the .NET ThreadPool class
 - At least 1 thread per CPU core
 - Tasks (work items) queue on the threads (execution contexts)
 - “Millions” of tasks may exist for “a few threads”

TASK SCHEDULING – FIRST APPROACH



Many cores -> fine-grained tasks -> contention problems on work queue

THE DEFAULT TASK SCHEDULER



INLINING EXECUTION OF PENDING TASKS

- The default task scheduler may *inline* waiting tasks
- When task1 awaits task2, and task2 is not started at the time task1 awaits it, the scheduler can execute task2 immediately on task1's thread
- Happens only when...
 - task1 calls `Task.Wait(task2)` or `Task.WaitAll()`, and
 - The local work queue of the thread at which task1 is executing also contains task2

PASSING DATA USING CLOSURES

- Closures are easy!

```
public void DoWork() {  
    int data1 = 42;  
    string data2 = "The Answer to the Ultimate Question of " +  
        "Life, the Universe, and Everything";  
  
    Task.Run(()=>  
    {  
        Console.WriteLine(data2 + ": " + data1);  
    });  
}
```

Delegate refers to state outside its scope (data1 and data2), so compiler creates a "closure" for data1 and data2 to make them accessible to the delegate

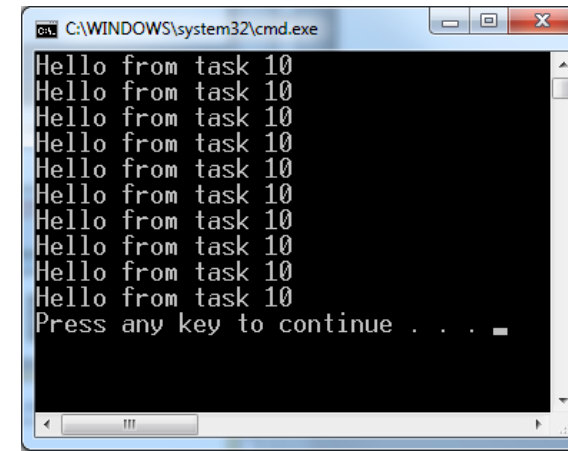
PASSING DATA USING CLOSURES

- Closures are also error-prone!
 - Copy the variable to be captured to a local variable before capturing it.

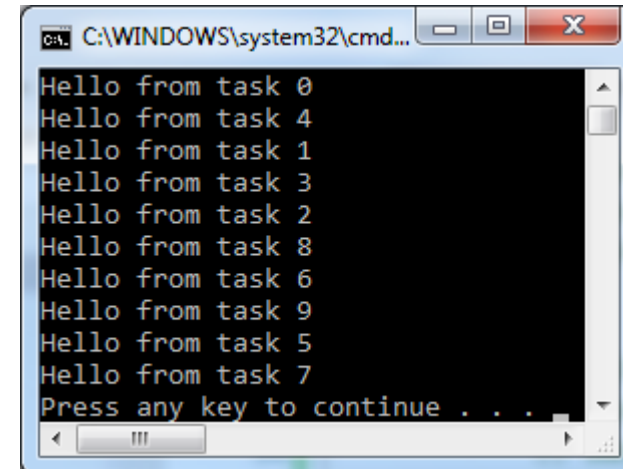
```
for(int i=0; i< 10; i++)  
    Task.Run(()=> {  
        Console.WriteLine(" Hello from task " + i);  
    });
```

Capture value of *i* in
local variable *localI*

```
for(int i=0; i< 10; i++)  
{  
    var localI = i;  
    Task.Run(()=> {  
        Console.WriteLine(" Hello from task " + localI);  
    });  
}
```



```
C:\WINDOWS\system32\cmd.exe  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Hello from task 10  
Press any key to continue . . .
```



```
C:\WINDOWS\system32\cmd...  
Hello from task 0  
Hello from task 4  
Hello from task 1  
Hello from task 3  
Hello from task 2  
Hello from task 8  
Hello from task 6  
Hello from task 9  
Hello from task 5  
Hello from task 7  
Press any key to continue . . .
```

PASSING DATA USING STATE OBJECTS

- Passing objects containing state to task:

```
int taskNo = 42;  
  
Task.Factory.StartNew((state) =>  
{  
    Console.WriteLine(" Hello from task " + (int) state);  
}, taskNo);
```

taskNo is the value of state when task is run.
taskNo is captured properly
Passing objects can *not* be done with Task.Run()!

PASSING DATA USING STATE OBJECTS

- Passing objects containing state *and* methods:

```
class Work
{
    public int Data1;
    public string Data2;
    public void Run() {
        Console.WriteLine(Data1 + ": " + Data2);
    }
}

public static void Main()
{
    Work w = new Work();
    w.Data1 = 42;
    w.Data2 = "The Answer to the Ultimate Question of...";
    Task.Run(w.Run);
}
```

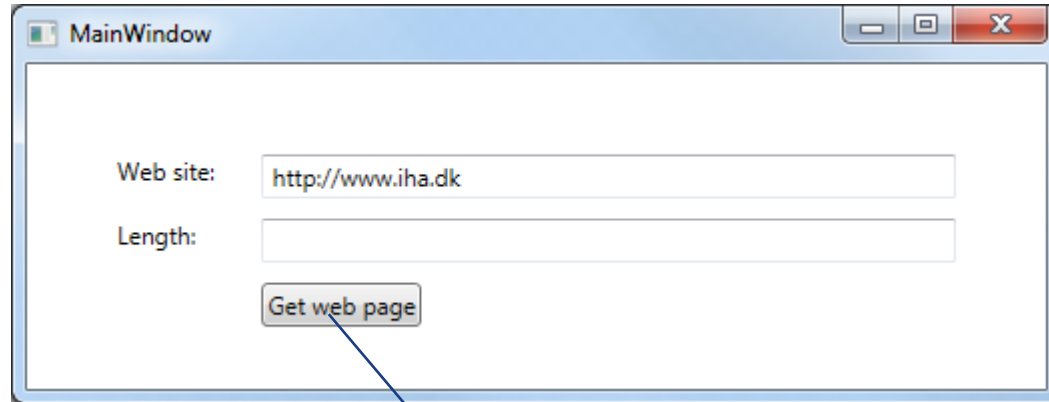
<https://learn.microsoft.com/en-us/dotnet/api/system.action?view=net-6.0>



TASK-BASED ASYNCHRONOUS PATTERN (TAP)

- .Net 4.5 exposed asynchronous versions of many operations following the Task-based Asynchronous Pattern (TAP)
- TAP is used with the `async` modifier and can include `await` expressions
- Allows calling thread (e.g. UI thread) to remain responsive while heavy operations execute

TAP



`GetStringAsync()` actually returns a `Task<string>`, this is unwrapped by `await`

```
private async void btnGetHtml_Click(object sender, RoutedEventArgs e)
{
    tbxLength.Text = "Fetching...";
    string url = tbxUrl.Text;
    HttpClient client = new HttpClient();
    string text = await client.GetStringAsync(url);
    tbxLength.Text = text.Length.ToString();
}
```

The boxed code is wrapped as a new Task which will be scheduled on the same thread when the operation returns

TAP - AWAIT

- `await` allows you to await the completion of an asynchronous operation
 - The rest of the code (from `await` onwards) is executed when the awaited operation returns.
 - The awaited operation is executed in the hardware, the drivers, and perhaps another thread
 - Use it for time-consuming IO (Input/Output)
- Awaiting is accomplished *without* blocking the calling thread (often the UI thread)
 - Instead, the awaited method returns immediately after `await`
 - If a result is immediately available, the caller continues
 - If not, a continuation is scheduled to execute on the caller thread when the awaited operation has been completed

ConfigureAwait

- `ConfigureAwait(false)`
 - Avoid queuing the task callback
 - Code after 'await' does not necessarily run in the same context
 - Improves performance
- `ConfigureAwait(true)`
 - == default behavior



AARHUS
UNIVERSITY